

# **EE5179: DEEP LEARNING IN IMAGING**

## **ASSIGNMENT 4: AUTOENCODERS REPORT**

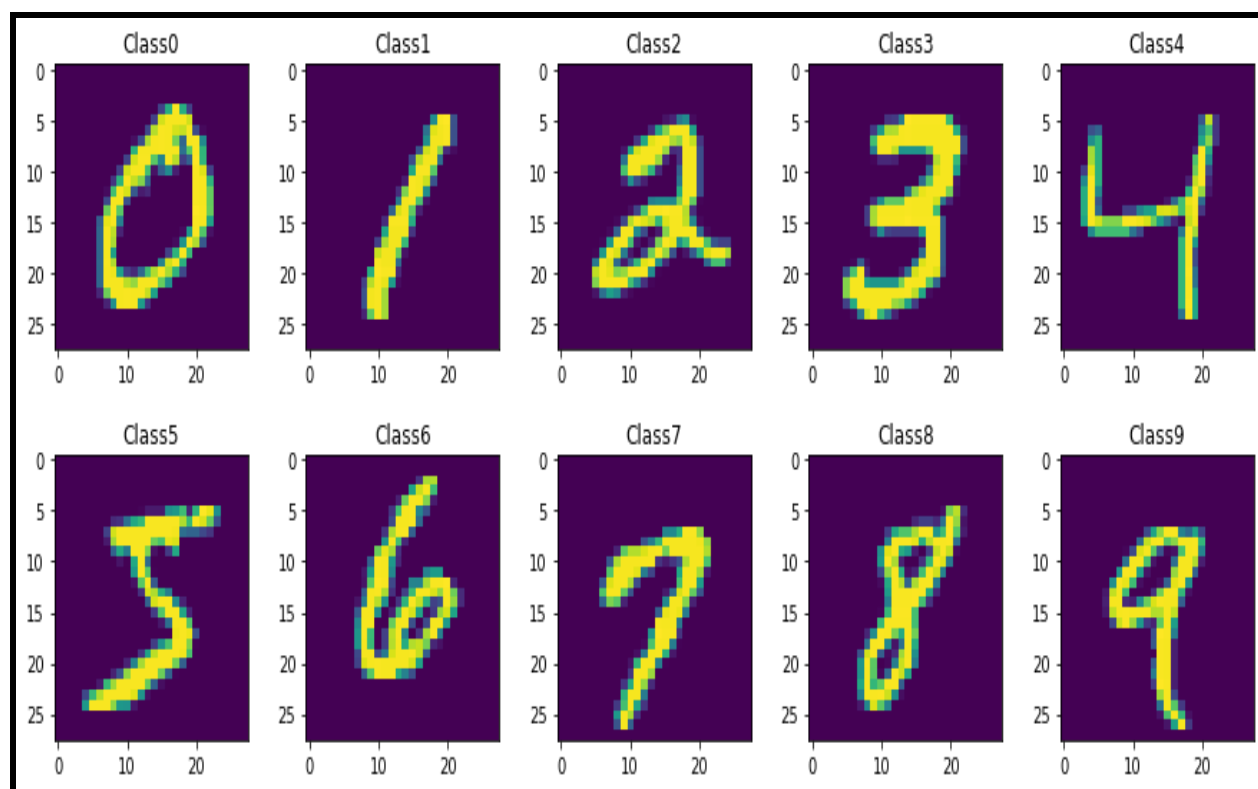
**RAVEENA RAI (ED21S006)**

## Resources used:

1. Google Colab
2. MNIST Dataset of handwritten digits
3. Tutorials provided in the class of the course 'Deep Learning for Imaging'
4. Online references.

## About the dataset:

1. The dataset used in the algorithm is the MNIST Digits Recognition dataset.
2. It comprises handwritten digits, pre-processed to ensure that the digits are centered and size normalized.
3. The train set consists of 60,000 images, and the test set consists of 10,000 images.
4. The training dataset was divided into two sets: training (50000 images) and validation(10000 images).
5. The train data has been normalised before using it in the assignment.
6. The images have been flatten ( $28*28=784$ ).



## Q-1: COMPARING PCA AND AUTOENCODER:

### A) PCA:

1. PCA was performed on the MNIST dataset and first 30 eigenvalues with their corresponding eigenvectors were used.
2. To observe the performance of the method, reconstruction was performed.

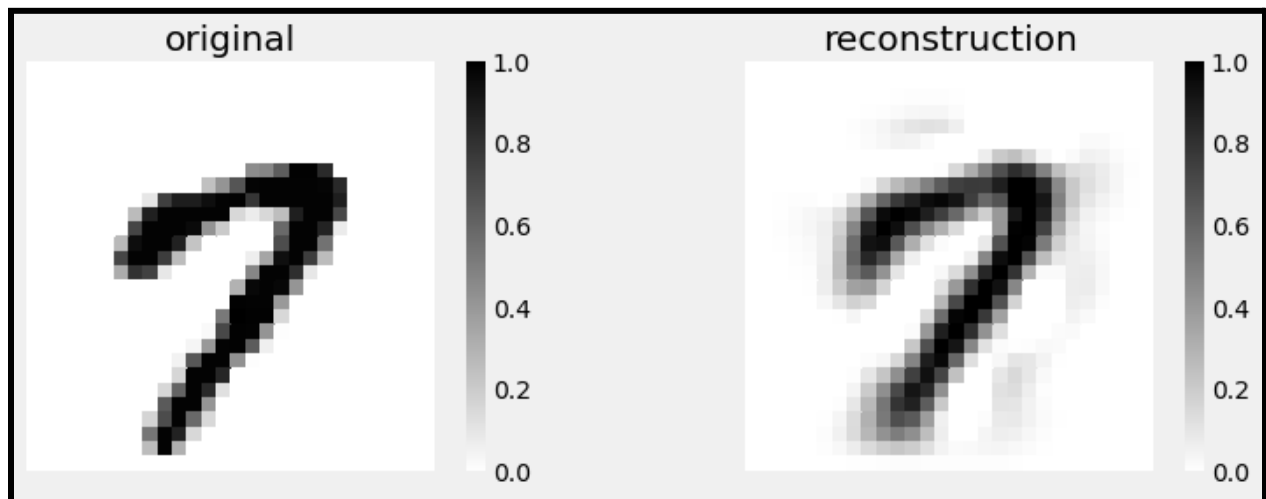
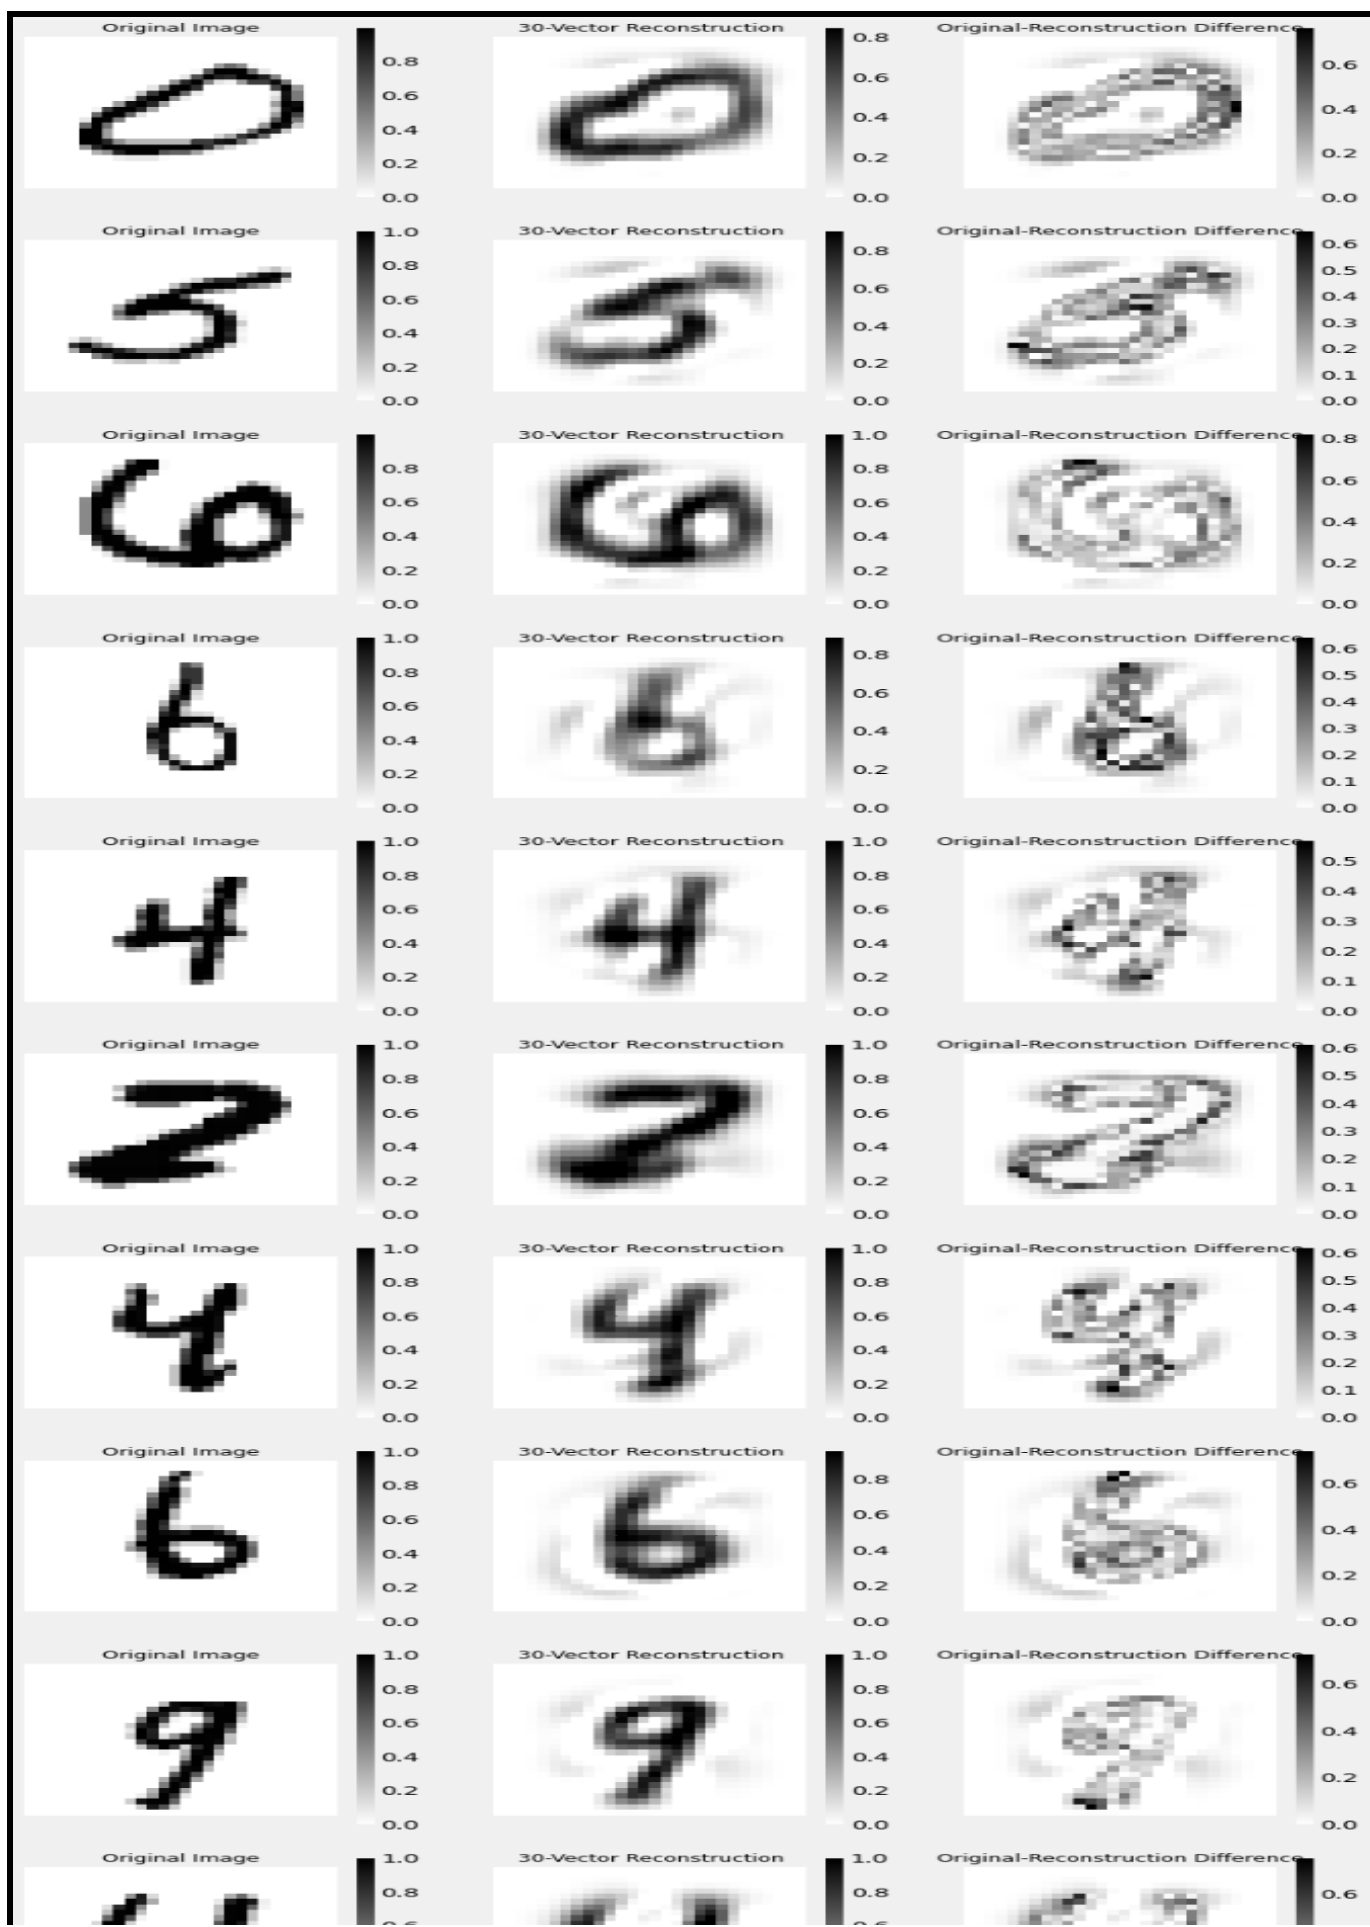


Fig: reconstruction



1. The figure shows the difference as obtained, between the original image and reconstructed image.
2. It can be seen that there is a lot of difference observed in the case of digits which have holes like 6, 0, 9, etc.

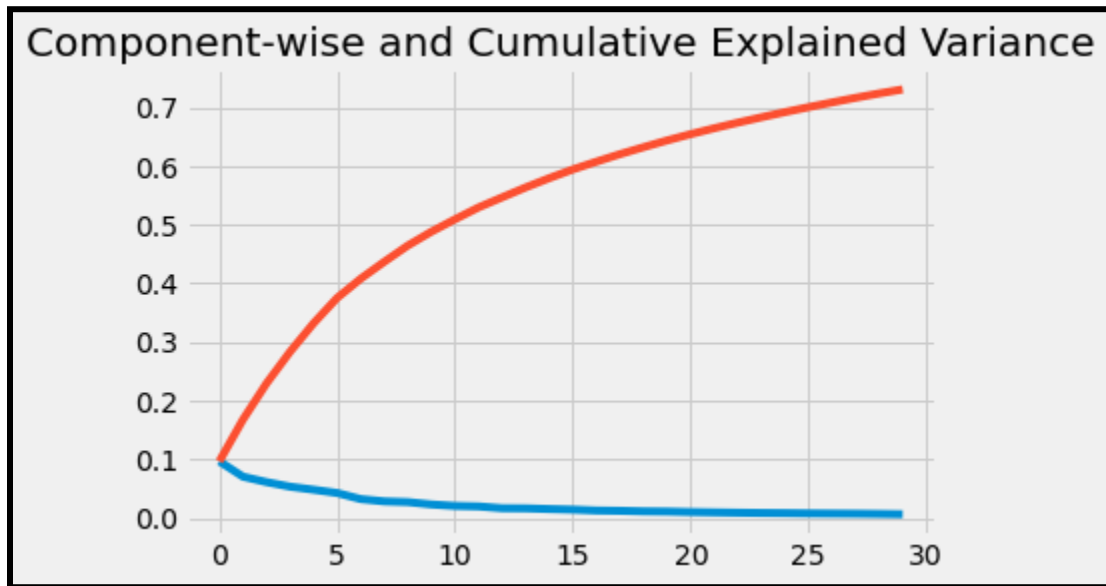


Fig: component- wise and cumulative explained variance

3. Looking at this plot, it seems like the first and second components add a lot of information about the variables to the model, but every component after that is not particularly worthwhile.
4. If we were using PCA as a processing step ourselves, we might want to cut off at  $n=2$ .

2nd Percentile	10th Percentile	25th Percentile	50th Percentile	75th Percentile	90th Percentile	98th Percentile
1	4	1	5	8	3	2
4	4	9	6	0	5	2
1	1	7	3	2	9	2
5	2	1	2	1	8	5
1	2	7	9	3	4	0
1	2	5	9	8	2	0
1	1	1	9	2	8	6
1	4	1	6	2	2	8

fig: Returns the data, which is the q-quartile fit for the given vector.

5. The 0th, 25th, 50th, 75th, and 100th percentile of the dataset. Records (quantiles) very close to the edges of the usefulness of a variable are often very informative as to what that variable "does".
6. This is true of ordinary variables, and all the more true of our computed ones, which we need all the help we can get interpreting.



## Autoencoders:

1. A simple autoencoder was constructed as follows:

Model: "autoencoder"

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 28, 28, 1)]	0
flatten_1 (Flatten)	(None, 784)	0
dense_7 (Dense)	(None, 512)	401920
dense_8 (Dense)	(None, 256)	131328
dense_9 (Dense)	(None, 128)	32896
dense_10 (Dense)	(None, 30)	3870
dense_11 (Dense)	(None, 128)	3968
dense_12 (Dense)	(None, 256)	33024
dense_13 (Dense)	(None, 784)	201488
reshape_1 (Reshape)	(None, 28, 28, 1)	0

Total params: 808,494

Trainable params: 808,494

Non-trainable params: 0

2. reLu was used as an activation function.
3. The train and validation plots were obtained.
4. Batch size was selected arbitrarily as 32, and the model was trained for 10 epochs with mse loss and learning rate as 0.001 (arbitrary).
5. validation accuracy obtained was 81.43 %.




Fig: training accuracy




Fig: validation accuracy




Fig: training loss

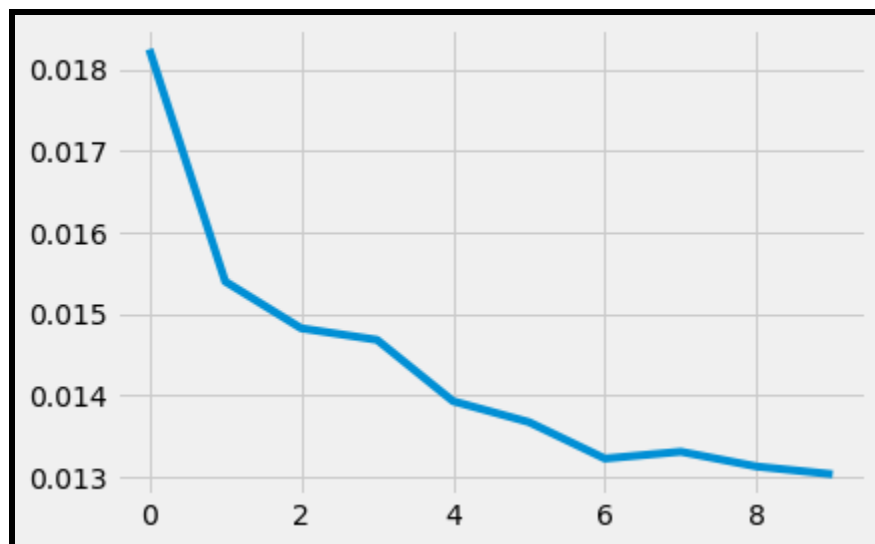


Fig: validation loss




Fig: reconstruction in PCA

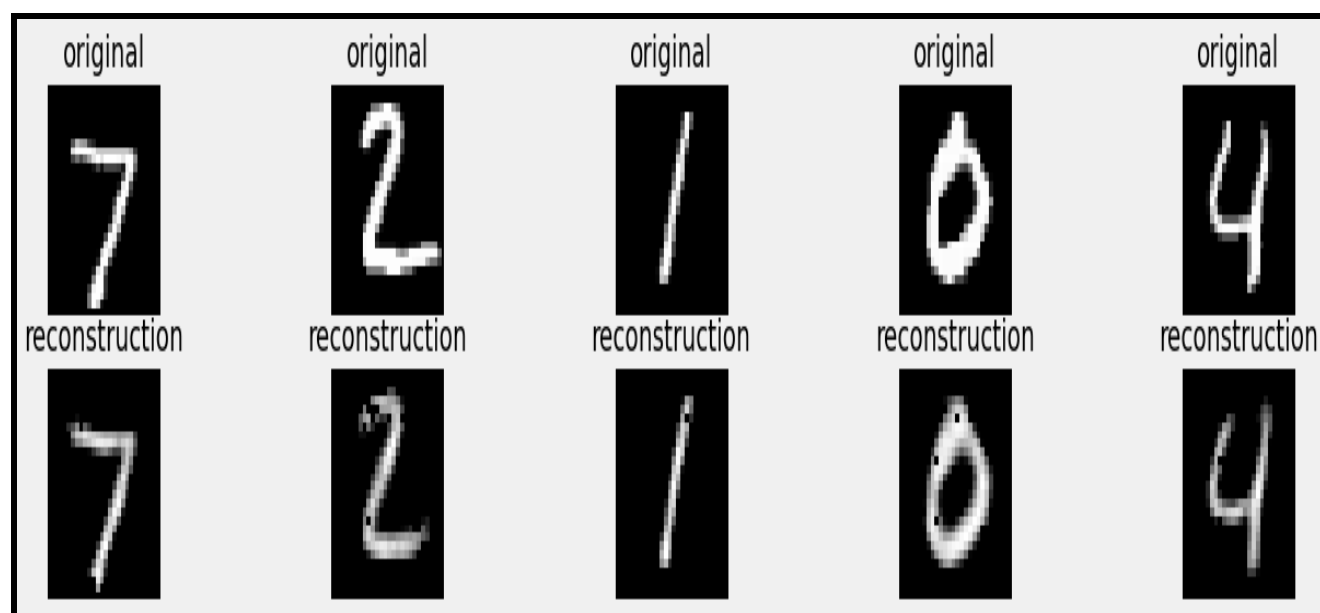


Fig: reconstruction in autoencoder

## Q-2: Experimenting with hidden units of varying sizes

1. Training a standard auto-encoder with the following architecture:

- $fc(x)-fc(784)$

Here,  $x$  is the size of the hidden unit. The architecture consists of only a hidden layer and the output layer. Using  $x = [64, 128, 256]$

1.  $fc(64)-fc(784)$

```
Model: "autoencoder"
```

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 28, 28, 1)]	0
flatten_4 (Flatten)	(None, 784)	0
dense_20 (Dense)	(None, 64)	50240
dense_21 (Dense)	(None, 64)	4160
dense_22 (Dense)	(None, 784)	50960
reshape_4 (Reshape)	(None, 28, 28, 1)	0

```
=====  
Total params: 105,360  
Trainable params: 105,360  
Non-trainable params: 0  
=====
```

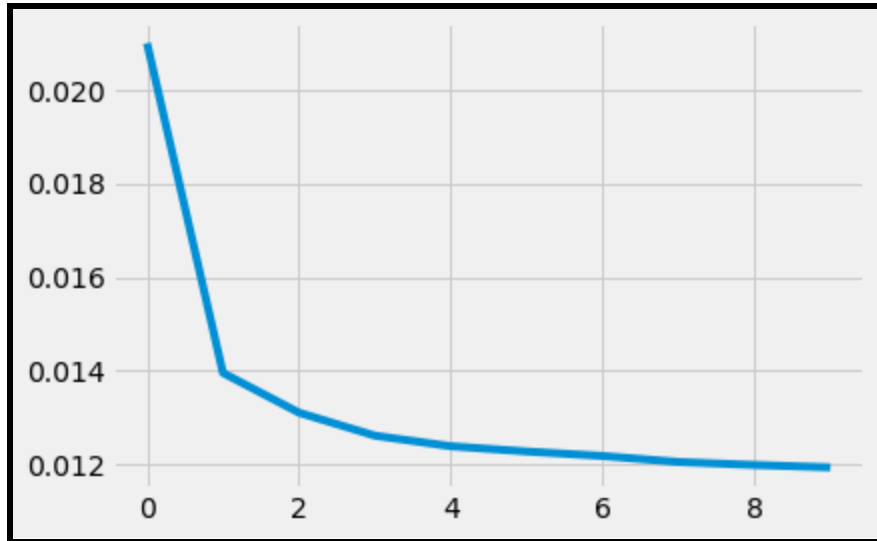


Fig: training loss




Fig:validation loss




Fig: training accuracy




Fig: validation accuracy

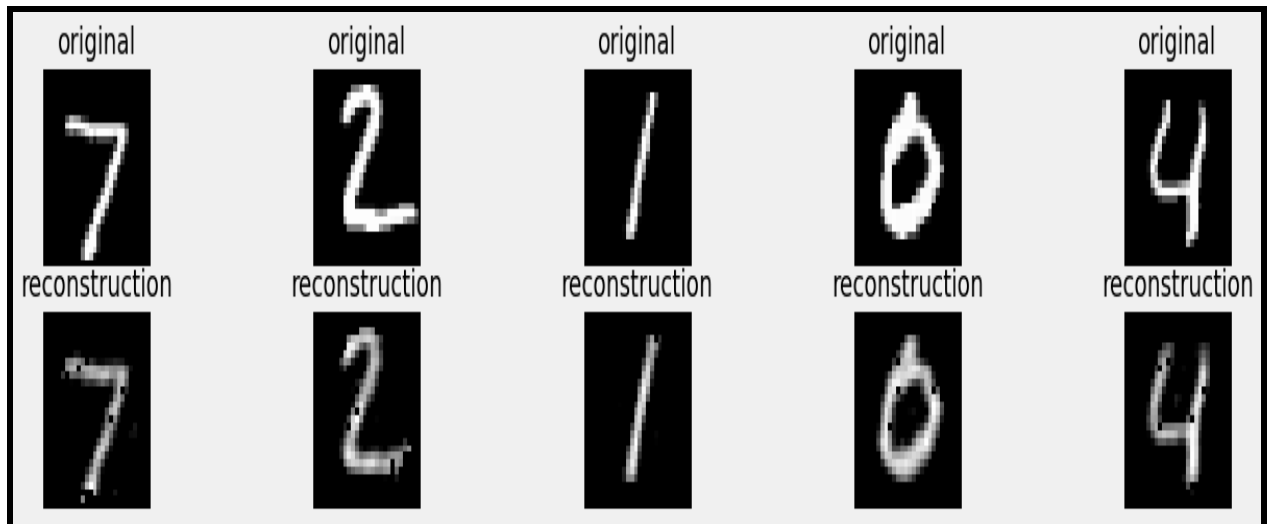


Fig: reconstruction

fc(128)-fc(784)

Model: "autoencoder"

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 28, 28, 1)]	0
flatten_5 (Flatten)	(None, 784)	0
dense_23 (Dense)	(None, 128)	100480
dense_24 (Dense)	(None, 128)	16512
dense_25 (Dense)	(None, 784)	101136
reshape_5 (Reshape)	(None, 28, 28, 1)	0

=====  
Total params: 218,128  
Trainable params: 218,128  
Non-trainable params: 0

Validation accuracy =81.53%







Fig:training loss



Validation loss




Fig: training accuracy


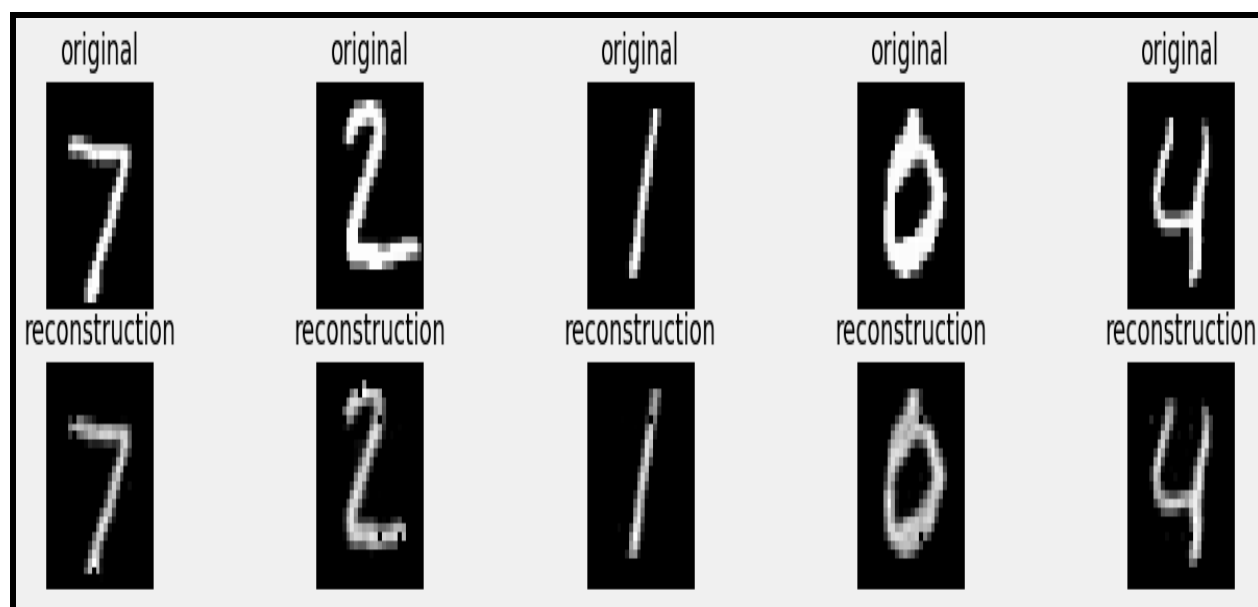


Fig: validation accuracy





fc(256)-fc(784)


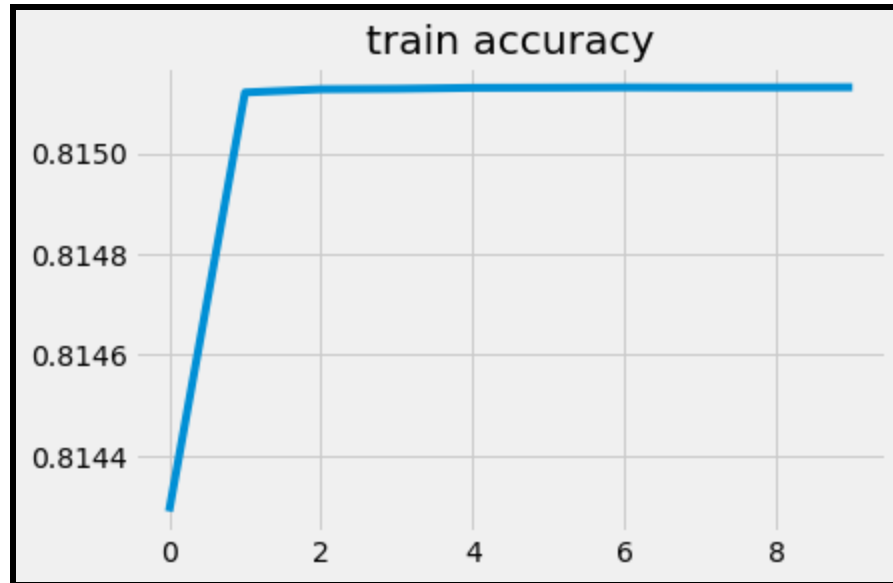
Model: "autoencoder"

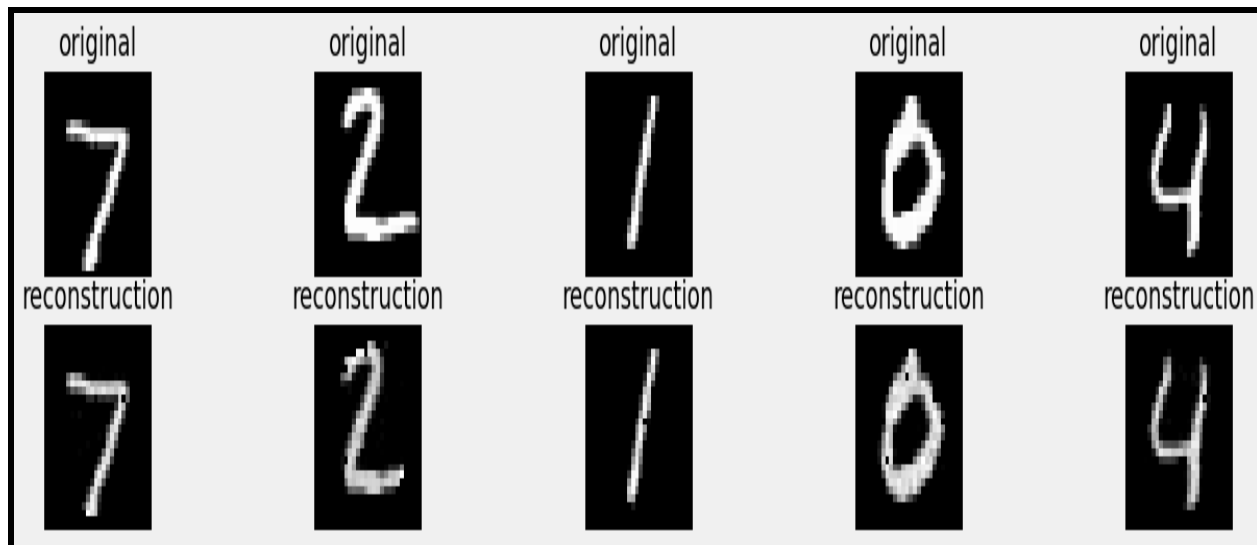
Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 28, 28, 1)]	0
flatten_6 (Flatten)	(None, 784)	0
dense_26 (Dense)	(None, 256)	200960
dense_27 (Dense)	(None, 256)	65792
dense_28 (Dense)	(None, 784)	201488
reshape_6 (Reshape)	(None, 28, 28, 1)	0

=====


Total params: 468,240  
Trainable params: 468,240  
Non-trainable params: 0







Comparison between the model with varying hidden sizes :






Fig: fc(64)




Fig: fc(128)




Fig: fc(256)

MNIST Fashion dataset:




fig : reconstruction of mnist fashion dataset images with trained autoencoder



Random noise images through trained autoencoder:


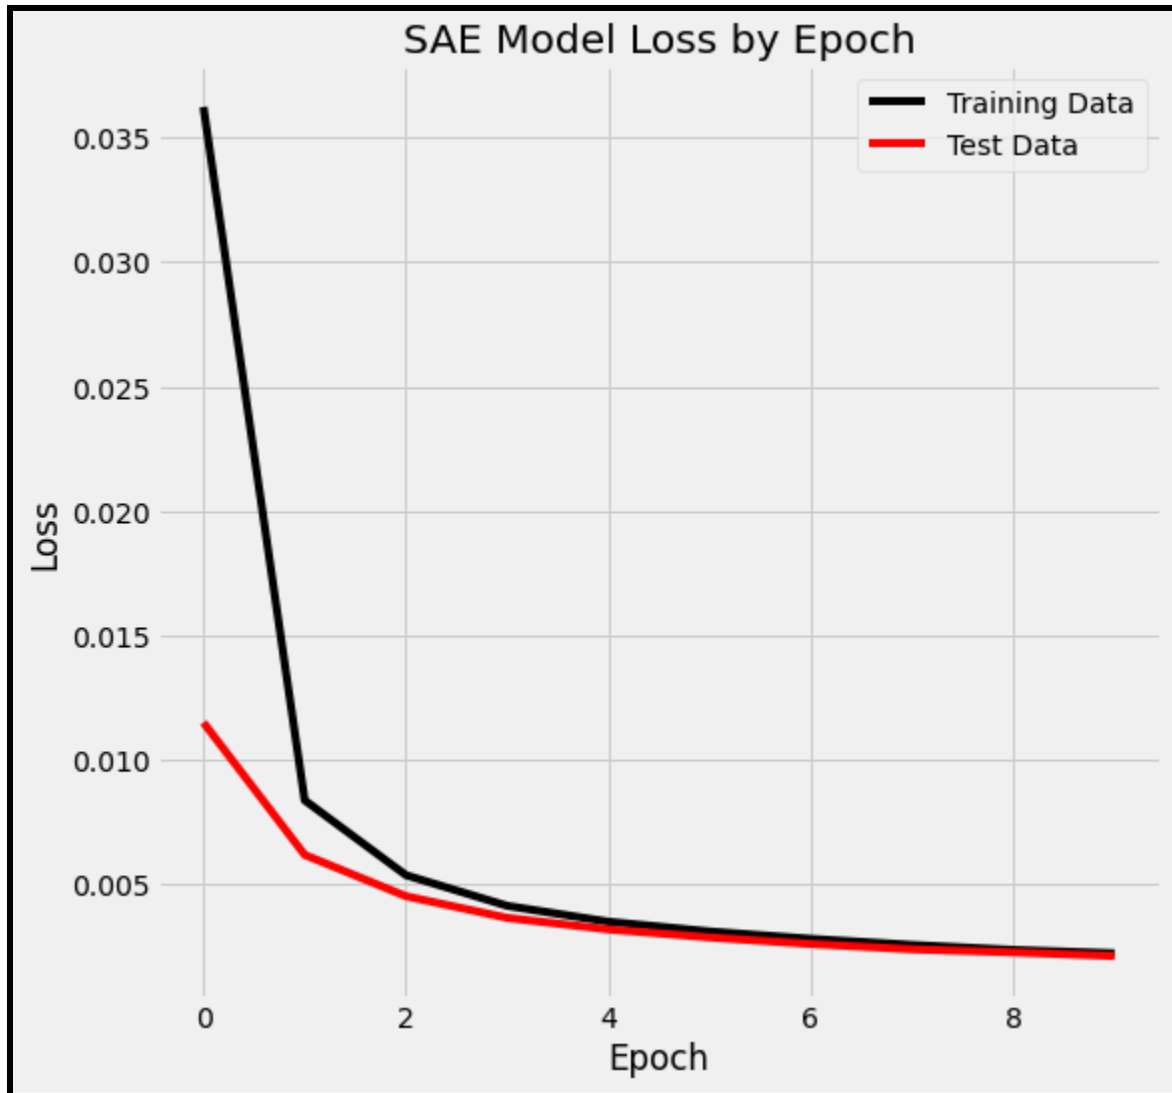


Fig:random noise image reconstruction using trained encoder.

### Q-3: Sparse autoencoders:

Trained an over-complete autoencoder with L1 sparsity regularization.

1. Sparsity value: 0.0001






Fig: accuracy plot





Fig: reconstruction

2. Sparsity value: 0.001






Fig: model accuracy






Fig: reconstruction

3. Sparsity value: 0.000001










Fig: reconstruction




Fig: filter for sparsity




Fig: filter for standard encoder

Observation:

1. In sparsity regulated autoencoder, many values are hidden compared to a standard autoencoder.

## Q-4: Denoising Autoencoder:

Designing denoising autoencoder with one hidden unit( 256).  
All models trained with batch size= 32 and epoch=10

Noise level=0.7

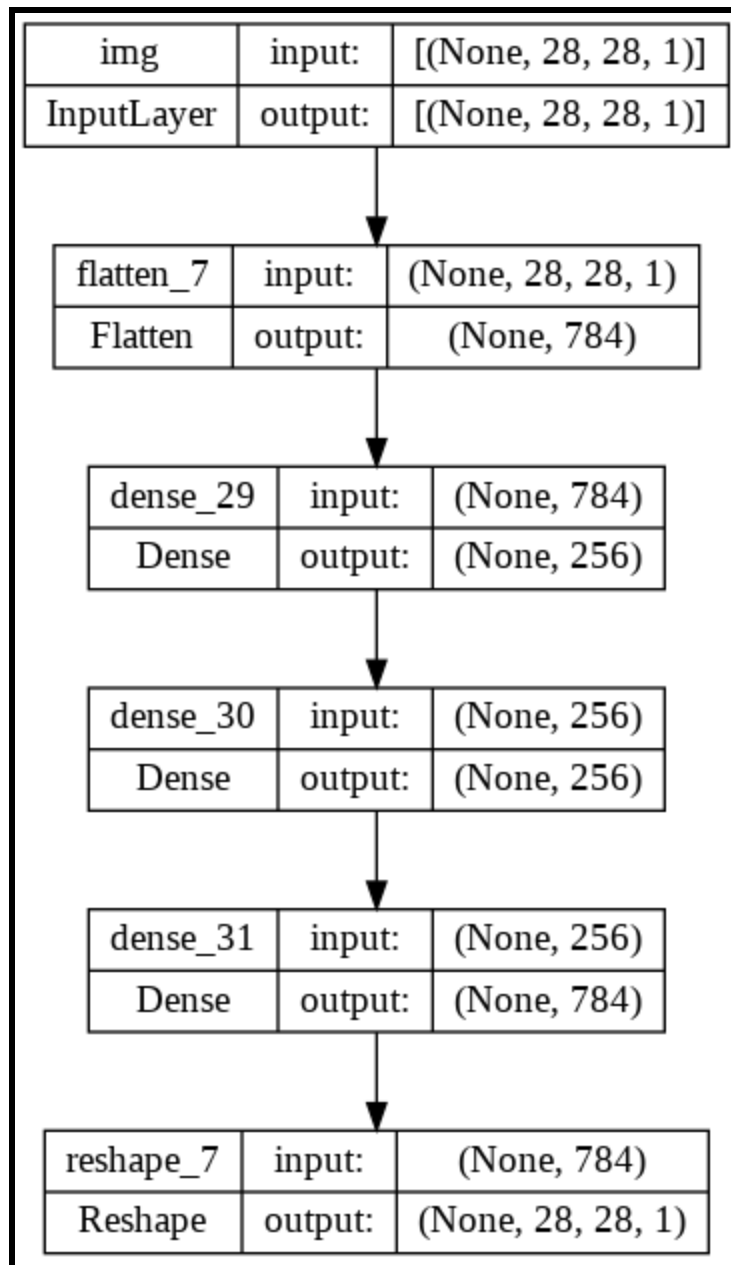


Fig: example images as generated in noisy data.

Model: "autoencoder"

Layer (type)	Output Shape	Param #
img (InputLayer)	[(None, 28, 28, 1)]	0
flatten_7 (Flatten)	(None, 784)	0
dense_29 (Dense)	(None, 256)	200960
dense_30 (Dense)	(None, 256)	65792
dense_31 (Dense)	(None, 784)	201488
reshape_7 (Reshape)	(None, 28, 28, 1)	0

=====  
Total params: 468,240  
Trainable params: 468,240  
Non-trainable params: 0



validation accuracy= 81.03%




Fig: original and reconstructed

Noise level: 0.3




Fig: noisy images




Fig: original and reconstructed

Noise level = 0.5



Fig: noisy images




Fig: original and reconstructed

Noise level = 0.8



Fig: noisy images




Fig: original and reconstructed

Noise level= 0.9

validation accuracy= 80.7%



Fig: noisy images




Fig: original and reconstructed

Q-5: convolutional autoencoder:



input_1	input:	[(None, 28, 28, 1)]
InputLayer	output:	[(None, 28, 28, 1)]



conv2d	input:	(None, 28, 28, 1)
Conv2D	output:	(None, 28, 28, 8)



max_pooling2d	input:	(None, 28, 28, 8)
MaxPooling2D	output:	(None, 14, 14, 8)



conv2d_1	input:	(None, 14, 14, 8)
Conv2D	output:	(None, 14, 14, 16)



max_pooling2d_1	input:	(None, 14, 14, 16)
MaxPooling2D	output:	(None, 7, 7, 16)



conv2d_2	input:	(None, 7, 7, 16)
Conv2D	output:	(None, 7, 7, 16)



max_pooling2d_2	input:	(None, 7, 7, 16)
MaxPooling2D	output:	(None, 4, 4, 16)



conv2d_3	input:	(None, 4, 4, 16)
Conv2D	output:	(None, 4, 4, 16)



up_sampling2d	input:	(None, 4, 4, 16)
UpSampling2D	output:	(None, 8, 8, 16)



conv2d_4	input:	(None, 8, 8, 16)
Conv2D	output:	(None, 8, 8, 16)



conv2d_5	input:	(None, 8, 8, 16)
----------	--------	------------------

Trained with epoch=10 and batch size = 32:

1. unpooling:

Validation accuracy: 81.4%

Error: 0.086




Fig: reconstructed images from test dataset

(Please note: the remaining part of the assignment will be submitted as a different report and code due to time shortage.)